



ĐẠI HỌC
BÁCH KHOA HÀ NỘI
HANOI UNIVERSITY
OF SCIENCE AND TECHNOLOGY

Real-time Weather Data Analysis and Forecasting System in Vietnam

Big Data Storage and Processing – IT4043E
Group 10

ONE LOVE. ONE FUTURE.



Table of contents

- I. Problem Definition
- II. Architecture and Design
- III. Implementation Details
- IV. Lesson learned



I. Problem Definition

- Vietnam's increasing climate variability and the limitations of traditional forecasting methods highlight the need for a real-time, scalable system that integrates high-frequency data to deliver accurate, localized weather predictions.
- The problem is quintessential for “Big Data”:

Volume

Historical weather data can reach hundreds of thousands of rows, requiring fast read and write operations.

Velocity

Weather data arrives in real time and must be ingested and processed within seconds to be useful.

Variety

Data comes from multiple sources and formats (JSON and CSV) and must be unified into a single structure.

Veracity

Sensor noise and inconsistent fields require real-time data cleaning and validation

Value

Accurate real-time weather forecasting delivers significant socio-economic benefits for Vietnam.

I. Problem Definition

Scope


- Geographical Coverage: Analysis of weather data across 34 cities in Vietnam
- Data Sources: Open-source weather data collected via APIs from **OpenWeatherAPI** and **Open-Meteo**
- Temporal Focus: Short-term forecasting using 24-hour input data to predict next-hour temperature (historical data stored on S3 for future extension)
- Output: An interactive dashboard displaying weather analytics and next-hour temperature predictions

Limitations

- Hardware Constraints: High-resolution weather modeling requires substantial CPU/GPU resources and costly cloud infrastructure
- Data Availability: Limited or missing sensor coverage in some regions affects data completeness
- Forecast Uncertainty: Atmospheric systems are inherently chaotic, limiting prediction accuracy over longer time horizons

II. Architecture and Design

- Lambda Architecture covering Ingestion, Processing, Storage, Visualization



II. Architecture and Design

- Data Ingestion Layer



OpenMeteo: Historical weather data



OpenWeather Current weather data

```
{
  "latitude": 21.0,
  "longitude": 105.875,
  "generationtime_ms": 0.1646280288696289,
  "utc_offset_seconds": 0,
  "timezone": "GMT",
  "timezone_abbreviation": "GMT",
  "elevation": 19.0,
  "hourly_units": {
    "time": "iso8601",
    "temperature_2m": "°C",
    "relative_humidity_2m": "%",
    "surface_pressure": "hPa",
    "cloud_cover": "%",
    "wind_speed_10m": "km/h",
    "wind_direction_10m": "°",
    "weather_code": "wmo code"
  },
  "hourly": {
    "time": [
      "2025-12-27T00:00", "2025-12-27T01:00", "2025-12-27T02:00",
      ...
    ],
    "temperature_2m": [
      15.0, 15.4, 16.1, ...
    ],
    ...
  }
}
```

```
[2025-12-28 13:00:35.980135] Collecting weather data for Hanoi...
Data from OPENWEATHER:
{
  'base': 'stations',
  'clouds': {'all': 100},
  'cod': 200,
  'coord': {'lat': 21.0245, 'lon': 105.8412},
  'dt': 1766901228,
  'id': 1581130,
  'main': {
    'feels_like': 19.54,
    'grnd_level': 1017,
    'humidity': 57,
    'pressure': 1018,
    'sea_level': 1018,
    'temp': 20,
    'temp_max': 20,
    'temp_min': 20
  },
  'name': 'Hanoi',
  'sys': {
    'country': 'VN',
    'id': 9308,
    'sunrise': 1766878327,
    'sunset': 1766917427,
    'type': 1
  },
  'timezone': 25200,
  'visibility': 10000,
  'weather': [{
    'description': 'overcast clouds',
    ...
  }]
}
```

II. Architecture and Design

- Data Ingestion Layer

Data sources



Approach Tried

- Removed unnecessary columns to reduce data size
- Kept only features needed for ML training and real-time weather display

Final Solution

Dropped unused columns to significantly speed up ML processing

Key Takeaways

- Different APIs have inconsistent data structures and require preprocessing
- Flattening JSON simplifies data for efficient processing
- Reducing data dimensions improves performance in real-time pipelines

```
[3]:
```


	city	timestamp	description	temp	humidity	pressure	wind_speed	wind_deg	cloudiness	feels_like	temp_min	temp_max	win
0	Hanoi	2025-11-29 07:00:00	clear sky	23.5	50	1011.3	4.00	360	0	NaN	NaN	NaN	
1	Hanoi	2025-11-29 08:00:00	clear sky	23.3	49	1010.4	2.30	39	0	NaN	NaN	NaN	
2	Hanoi	2025-11-29 09:00:00	clear sky	22.7	51	1010.8	2.50	98	0	NaN	NaN	NaN	
3	Hanoi	2025-11-29 10:00:00	clear sky	21.0	62	1010.9	1.80	127	0	NaN	NaN	NaN	
4	Hanoi	2025-11-29 11:00:00	clear sky	19.0	62	1011.6	2.30	162	0	NaN	NaN	NaN	
5	Hanoi	2025-11-29 12:00:00	clear sky	18.3	62	1011.4	2.20	171	0	NaN	NaN	NaN	
6	Hanoi	2025-11-29 13:00:00	clear sky	17.3	65	1011.9	3.20	153	0	NaN	NaN	NaN	
7	Hanoi	2025-11-29 14:00:00	clear sky	16.5	68	1012.0	6.40	133	0	NaN	NaN	NaN	
8	Hanoi	2025-11-29 15:00:00	clear sky	14.7	76	1012.5	4.70	122	0	NaN	NaN	NaN	
9	Hanoi	2025-11-29 16:00:00	clear sky	13.8	79	1012.3	3.90	124	0	NaN	NaN	NaN	

II. Architecture and Design

Stream Processing

Apache Kafka

- Kafka handles stream data. In this case, stream data is the current weather data from OpenWeather API.




```
{"city": "Gia Lai", "city_id": 1581088, "country": "VN", "event_time": "2025-12-28T05:00:00+00:00", "dt": 1766897730, "sunrise": 1766876925, "sunset": 1766917672, "timezone": 25200, "coord_lat": 13.75, "coord_lon": 108.25, "weather_id": 804, "weather_main": "Clouds", "weather_description": "overcast clouds", "weather_icon": "04d", "temp": 25.33, "feels_like": 25.35, "temp_min": 25.33, "temp_max": 25.33, "pressure": 1014, "humidity": 55, "sea_level": 1014, "grnd_level": 963, "visibility": 10000, "wind_speed": 4.16, "wind_deg": 72, "wind_gust": 5.27, "cloudiness": 91}
{"city": "Hanoi", "city_id": 1581130, "country": "VN", "event_time": "2025-12-28T05:00:00+00:00", "dt": 1766898022, "sunrise": 1766878327, "sunset": 1766917427, "timezone": 25200, "coord_lat": 21.0245, "coord_lon": 105.8412, "weather_id": 804, "weather_main": "Clouds", "weather_description": "overcast clouds", "weather_icon": "04d", "temp": 19, "feels_like": 18.44, "temp_min": 19, "temp_max": 19, "pressure": 1019, "humidity": 57, "sea_level": 1019, "grnd_level": 1018, "visibility": 10000, "wind_speed": 1.36, "wind_deg": 349, "wind_gust": 1.38, "cloudiness": 100}
```


II. Architecture and Design

Stream Processing

Apache Spark (Stream Spark + Bronze)

- Real-time weather data ingested from **OpenWeather API** via **Kafka**
- **Spark Structured Streaming** processes event-time data continuously
- Raw events stored in **Bronze (Parquet)** for durability
- **Windowed aggregation + watermark** handle late data
- Aggregated results **upsurged** to **Elasticsearch** for real-time analytics




II. Architecture and Design

Batch Processing

Amazon S3

- Amazon S3 serves as a storage for historical data
- Historical weather data stored in **Amazon S3**
- **Spark batch jobs** process 24-hour data windows
- **CronJob runs every 45 minutes** to add datasets
- Processed data written to **Silver layer** and **ML-ready tables**



```
download: s3://hust-buck  
city: string  
timestamp: timestamp[us]  
description: string  
temp: double  
humidity: double  
pressure: double  
wind_speed: double  
wind_deg: double  
cloudiness: double
```

II. Architecture and Design

Batch Processing

Time-Series Data Storage Evaluation

- Large-scale time-series weather data requires efficient and scalable storage.
- Multiple storage formats were evaluated for ingestion and analytics performance



- Row-based storage format.
- Excellent for fast ingestion and schema evolution.
- Inefficient for analytics since entire rows must be read even when querying a single column.



- Columnar format optimized for Hive.
- Highest compression ratio among tested formats.
- Less compatible with Python and broader data science workflows outside Hadoop.



- Columnar format designed for analytical workloads and nested data.
- Well-suited for transformed datasets in Silver and Gold layers.
- Strong ecosystem support across Spark and Python tools.

II. Architecture and Design

Batch Processing

Parquet chosen as the primary storage format

- Best balance of performance, usability, and ecosystem compatibility.
- Easier to work with than Avro and ORC for analytics and ML pipelines.



Objects (5)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

☐	Name	Type	Last modified	Size	Storage class
☐	Can Tho.parquet	parquet	November 30, 2025, 14:37:34 (UTC+07:00)	9.3 KB	Standard
☐	Da Nang.parquet	parquet	November 30, 2025, 14:37:31 (UTC+07:00)	9.1 KB	Standard
☐	Haiphong.parquet	parquet	November 30, 2025, 14:37:32 (UTC+07:00)	9.3 KB	Standard
☐	Hanoi.parquet	parquet	November 30, 2025, 14:37:28 (UTC+07:00)	9.3 KB	Standard
☐	Ho Chi Minh City.parquet	parquet	November 30, 2025, 14:37:29 (UTC+07:00)	9.3 KB	Standard


```
-rw-r--r-- 1 185 185 7.1K Dec 25 16:06 part-00002-fc21c3c9-a99d-4957-bd9f-012c45e7cc66-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 27 16:58 part-00002-fc4bba53-30b9-413e-88c1-d062d3d03632-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 28 10:15 part-00002-fc7d6cc7-79ad-46aa-9a67-1b15e9e0561a-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 27 17:10 part-00002-fccaae82-7698-4a1d-abcd-505e5ea029fd-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.4K Dec 28 09:55 part-00002-fd139514-bd05-48df-b3dc-b0f5dec812fa-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 28 03:32 part-00002-fd37cedf-c51e-47c5-86e8-f0f0f66f2af7-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.4K Dec 27 04:52 part-00002-fd3f78e2-e1d8-47de-94d0-7d6f9e71da25-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 26 03:59 part-00002-fd64a450-00ca-4e4e-bfd5-cac835d0f05a-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.4K Dec 27 04:23 part-00002-fd6b9ce4-5e36-4774-83e7-c186cdadc56f-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 25 10:37 part-00002-fda495c2-0f35-49d9-9a34-5e244fe37496-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 26 08:32 part-00002-fdabd5e7-776a-482e-91e7-05eaf6e4b21e-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 27 03:36 part-00002-fddef2b4-fb6f-446b-a78d-b690064f111c-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.5K Dec 28 00:46 part-00002-fde0eb17-7b0a-466f-81a3-ad6ff90286a0-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 27 02:59 part-00002-fe6919f6-5603-45c1-b3c9-0a2052359ccd-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 28 03:29 part-00002-fee5b2ca-c5cd-4013-b71d-3a7aaf19006e-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.5K Dec 27 23:12 part-00002-fe541ca-b64a-434e-a75a-14a5308b94d5-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 25 09:31 part-00002-fef93047-68da-4dab-99b0-6c0418370d99-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 26 03:00 part-00002-ff097666-8457-49da-a798-4715efaac432-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 26 04:14 part-00002-ff3b60fb-6bfc-4923-8a52-643836f7ceae-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 28 00:46 part-00002-ff9a8f47-2e93-46d5-b42a-8db9865e009f-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 27 16:57 part-00002-ffae2eb2-0f42-4548-b21e-997b31e4683b-c000.snappy.parquet
-rw-r--r-- 1 185 185 7.1K Dec 27 23:19 part-00002-ffc9a179-10b3-404d-ade3-72cbc98353a5-c000.snappy.parquet
```

II. Architecture and Design

Batch Processing

Silver data

- Read historical data stored in Amazon S3
- Spark batch job runs hourly (CronJob)
- Cleans, deduplicates, and add average temperature in 24 hours
- Outputs **Silver data** for **Gold data**





II. Architecture and Design

Batch Processing

Gold data

- Spark batch job runs every 1h10m (CronJob)
- Aggregates **Silver data** → hourly city-level metrics
- Computes alerts & derived features
- Writes **Gold data** to Elasticsearch for dashboards





II. Architecture and Design

Machine Learning model Implementation

Spark MLlib

- Hourly batch inference with Spark MLlib
- Uses last 24h S3 + recent Bronze data
- Predicts next-hour temperature per city
- Stores predictions in Elasticsearch for visualization





II. Architecture and Design

Visualization

Kibana

- Monitoring & analytics
- Gold + ML outputs
- Interactive dashboards



II. Architecture and Design

Deployment

Kubernetes

- Kubernetes provides a **unified orchestration layer** to run streaming, batch, and ML workloads on the same infrastructure.
- It handles **resource scheduling and isolation**, allowing Spark, Kafka, and Elasticsearch to coexist without hard coupling.
- Kubernetes enables **fault recovery and restart** of long-running services such as Kafka brokers and Spark streaming jobs.
- It simplifies **deployment and lifecycle management** through declarative manifests and operators (Strimzi, Spark Operator, Elastic Operator).

```
NAME: data-platform-control-plane
STATUS: Ready
ROLES: control-plane
AGE: 9d
VERSION: v1.34.0
INTERNAL-IP: 172.19.0.2
EXTERNAL-IP: <none>
OS-IMAGE: Debian GNU/Linux 12 (bookworm)
KERNEL-VERSION: 6.10.14-linuxkit
CONTAINER: container

Name: data-platform-control-plane
Labels: control-plane, beta.kubernetes.io/arch=amd64, beta.kubernetes.io/os=linux, kubernetes.io/arch=amd64, kubernetes.io/os=linux, kubernetes.io/hostname=data-platform-control-plane
Annotations: node.alpha.kubernetes.io/ttl: 0, volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Thu, 18 Dec 2025 17:08:39 +0700
Taints: <none>
Unschedulable: false
Leases:
  HolderIdentity: data-platform-control-plane
  AcquireTime: unset
  RenewTime: Sun, 28 Dec 2025 16:32:42 +0700
Conditions:
  Type      Status      LastHeartbeatTime      LastTransitionTime      Reason      Message
  --      --      --      --      --      --
MemoryPressure  False      Sun, 28 Dec 2025 16:31:21 +0700      Sun, 28 Dec 2025 14:03:56 +0700      KubeletHasSufficientMemory      kubelet has sufficient memory availa
DiskPressure     False      Sun, 28 Dec 2025 16:31:21 +0700      Sun, 28 Dec 2025 14:03:56 +0700      KubeletHasNoDiskPressure        kubelet has no disk pressure
PIDPressure      False      Sun, 28 Dec 2025 16:31:21 +0700      Sun, 28 Dec 2025 14:03:56 +0700      KubeletHasSufficientPID          kubelet has sufficient PID availabl
Ready            True       Sun, 28 Dec 2025 16:31:21 +0700      Sun, 28 Dec 2025 14:16:57 +0700      KubeletReady                     kubelet is posting ready status

Addresses:
  InternalIP: 172.19.0.2
  Hostname: data-platform-control-plane
Capacity:
  cpu: 12
  ephemeral-storage: 1045761844Ki
  hugepages-1Gi: 0
  hugepages-2Mi: 0
  hugepages-32Mi: 0
  hugepages-64Ki: 0
  memory: 12235624Ki
  pods: 110
Allocatable:
  cpu: 12
  ephemeral-storage: 1045761844Ki
  hugepages-1Gi: 0
  hugepages-2Mi: 0
  hugepages-32Mi: 0
  hugepages-64Ki: 0
  memory: 12235624Ki
  pods: 110
System Info:
  Machine ID: 35ac27d5d5e14801861f7de38a71ba55
  System UUID: 35ac27d5d5e14801861f7de38a71ba55
  Boot ID: 802b052b-e139-4c43-9248-5dc5f98a8cc5
  Kernel Version: 6.10.14-linuxkit
  OS Image: Debian GNU/Linux 12 (bookworm)
  Operating System: linux
  Architecture: arm64
  Container Runtime Version: containerd://2.1.3
```



III. Implementation Details

Kafka Producer (Data sources readers): Weather-producer.yaml (for kafka)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: weather-kafka-producer
  namespace: data
spec:
  replicas: 1
  selector:
    matchLabels:
      app: weather-kafka-producer
  template:
    metadata:
      labels:
        app: weather-kafka-producer
    spec:
      containers:
        - name: producer
          image: weather-kafka-producer:latest
          imagePullPolicy: IfNotPresent
          env:
            - name: KAFKA_BROKER
              value: kafka-kafka-bootstrap.data.svc.cluster.local:9092
            - name: KAFKA_TOPIC
              value: weather-raw
            - name: OPENWEATHER_API_KEY
              valueFrom:
                secretKeyRef:
                  name: openweather-secret
                  key: api-key
            - name: POLL_INTERVAL
              value: "60"
```

Lifecycle

1. Deployment creates Pod
2. Pod starts Python script
3. Script:
 1. Fetches weather
 2. Produces messages to Kafka
4. If script crashes:
 1. Pod restarts
 2. Producer reconnects automatically

III. Implementation Details

S3

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: weather-history-to-s3
  namespace: data
spec:
  schedule: "*/45 * * * *" # mỗi 45 phút
  successfulJobsHistoryLimit: 2
  failedJobsHistoryLimit: 2

  jobTemplate:
    spec:
      backoffLimit: 2
      template:
        spec:
          restartPolicy: OnFailure

          containers:
            - name: s3-history
              image: weather-s3-history:latest
              imagePullPolicy: IfNotPresent

          env:
            - name: S3_BUCKET_NAME
              value: hust-bucket-storage

            - name: AWS_ACCESS_KEY_ID
              valueFrom:
                secretKeyRef:
                  name: aws-secret
                  key: access-key
```

III. Implementation Details

Bronze

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: bronze-to-es-batch-cron
  namespace: default
spec:
  schedule: "*/2 * * * *"
  concurrencyPolicy: Forbid
  successfulJobsHistoryLimit: 1
  failedJobsHistoryLimit: 3

  jobTemplate:
    spec:
      backoffLimit: 1
      template:
        spec:
          serviceAccountName: spark
          restartPolicy: Never

          containers:
            - name: submit-spark-job
              image: bitnami/kubectl:latest
              imagePullPolicy: IfNotPresent
              command:
                - /bin/sh
                - -c
                - |
                  cat <<'EOF' | kubectl apply -f -
                  apiVersion: sparkoperator.k8s.io/v1beta2
                  kind: SparkApplication
                  metadata:
```

Silver.yaml

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: weather-s3-to-silver-cron
  namespace: default
spec:
  schedule: "0 * * * *"
  concurrencyPolicy: Forbid
  successfulJobsHistoryLimit: 1
  failedJobsHistoryLimit: 3

  jobTemplate:
    spec:
      backoffLimit: 1
      template:
        spec:
          serviceAccountName: spark
          restartPolicy: Never

          containers:
            - name: submit-silver-spark-job
              image: bitnami/kubectl:latest
              imagePullPolicy: IfNotPresent
              command:
                - /bin/sh
                - -c
                - |
                  kubectl create -f - <<'EOF'
                  apiVersion: sparkoperator.k8s.io/v1beta2
                  kind: SparkApplication
                  metadata:
```

III. Implementation Details

Gold

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: silver-to-gold-es-cron
  namespace: default
spec:
  schedule: "10 * * * *" # chạy sau silver 10 phút
  concurrencyPolicy: Forbid
  successfulJobsHistoryLimit: 1
  failedJobsHistoryLimit: 3

  jobTemplate:
    spec:
      backoffLimit: 1
      template:
        spec:
          serviceAccountName: spark
          restartPolicy: Never

          containers:
            - name: submit-gold-spark-job
              image: bitnami/kubectrl:latest
              imagePullPolicy: IfNotPresent
              command:
                - /bin/sh
                - -c
                - |
                  kubectl create -f - <<'EOF'
                  apiVersion: sparkoperator.k8s.io/v1beta2
                  kind: SparkApplication
                  metadata:
```

Machine Learning Implementation

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: weather-ml-inference-cron
  namespace: default
spec:
  schedule: "0 * * * *" # ⌚ mỗi giờ
  concurrencyPolicy: Forbid
  successfulJobsHistoryLimit: 1
  failedJobsHistoryLimit: 3

  jobTemplate:
    spec:
      backoffLimit: 1
      template:
        spec:
          serviceAccountName: spark
          restartPolicy: Never

          containers:
            - name: submit-ml-spark-job
              image: bitnami/kubectrl:latest
              imagePullPolicy: IfNotPresent
              command:
                - /bin/sh
                - -c
                - |
                  kubectl create -f - <<'EOF'
                  apiVersion: sparkoperator.k8s.io/v1beta2
                  kind: SparkApplication
                  metadata:
```

III. Implementation Details

Elasticsearch

```
apiVersion: elasticsearch.k8s.elastic.co/v1
kind: Elasticsearch
metadata:
  name: es-weather
  namespace: observability
spec:
  version: 7.17.9
  http:
    tls:
      selfSignedCertificate:
        disabled: true
  nodeSets:
  - name: default
    count: 1
    config:
      node.store.allow_mmap: false
    podTemplate:
      spec:
        containers:
        - name: elasticsearch
          resources:
            requests:
              memory: 1Gi
              cpu: 250m
            limits:
              memory: 1Gi
              cpu: 1
```

Kibana

```
apiVersion: kibana.k8s.elastic.co/v1
kind: Kibana
metadata:
  name: kibana-weather
  namespace: observability
spec:
  version: 7.17.9
  count: 1

  elasticsearchRef:
    name: es-weather


  http:
    tls:
      selfSignedCertificate:
        disabled: true # để port-forward, debug

  podTemplate:
    spec:
      containers:
      - name: kibana
        resources:
          requests:
            cpu: 200m
            memory: 512Mi
          limits:
            cpu: 1
            memory: 1Gi
```

III. Implementation Details

Deployment strategy

We first deploy on docker using many Dockerfiles, with a docker-compose.yml file to run everything



```
spark-streamer:
  build:
    context: ./spark_streamer
  # volumes:
  #   - ./spark_stream.py:/app/spark_stream.py
  container_name: spark-streamer
  depends_on:
    - kafka
    - elasticsearch
  networks:
    - weather-net
  > Run Service
spark-batcher:
  build:
    context: ./spark_batcher      # Folder containing Dockerfile + script
  container_name: spark-batcher
  depends_on:
    - kafka
    - elasticsearch
  env_file:
    - .env
  networks:
    - weather-net
```

III. Implementation Details

Deployment strategy

Kubernetes

- Deployed on a **single-node Kubernetes cluster**, where one node acts as both **control plane and worker**.
- Uses **containerd** as the container runtime with CPU and memory sized for data processing workloads.
- Employs **PersistentVolumeClaims (PVCs)** for Spark checkpoints and intermediate data storage.
- Organizes services into **separate namespaces** to ensure logical isolation and easier management.

PodCIDR:	10.244.0.0/24					
PodCIDRs:	10.244.0.0/24					
ProviderID:	kind://docker/data-platform/data-platform-control-plane					
Non-terminated Pods:	(29 in total)					
Namespace	Name	CPU Requests	CPU Limits	Memory Requests	Memory Limits	Age
data	kafka-entity-operator-5968477844-cnhzk	0 (0%)	0 (0%)	0 (0%)	0 (0%)	17h
data	kafka-kafka-pool-0	0 (0%)	0 (0%)	0 (0%)	0 (0%)	109m
data	kafka-kafka-pool-1	0 (0%)	0 (0%)	0 (0%)	0 (0%)	113m
data	kafka-kafka-pool-2	0 (0%)	0 (0%)	0 (0%)	0 (0%)	113m
data	strimzi-cluster-operator-8457d7566-hx2cn	200m (1%)	1 (8%)	384Mi (3%)	384Mi (3%)	9d
data	weather-kafka-producer-5d7b968f59-f5c6p	0 (0%)	0 (0%)	0 (0%)	0 (0%)	3d
default	bronze-to-es-batch-97b6f89b644cf7c-exec-1	1 (8%)	0 (0%)	1433Mi (11%)	1433Mi (11%)	41s
default	bronze-to-es-batch-97b6f89b644cf7c-exec-2	1 (8%)	0 (0%)	1433Mi (11%)	1433Mi (11%)	41s
default	bronze-to-es-batch-driver	1 (8%)	0 (0%)	1433Mi (11%)	1433Mi (11%)	44s
default	checkpoint-debug	0 (0%)	0 (0%)	0 (0%)	0 (0%)	8d
default	silver-debug	0 (0%)	0 (0%)	0 (0%)	0 (0%)	8d
default	spark-debug	0 (0%)	0 (0%)	0 (0%)	0 (0%)	16h
default	weather-streaming-to-bronze-b311d79b638cfba4-exec-1	1 (8%)	0 (0%)	1433Mi (11%)	1433Mi (11%)	3h30m
default	weather-streaming-to-bronze-b311d79b638cfba4-exec-2	1 (8%)	0 (0%)	1433Mi (11%)	1433Mi (11%)	3h30m
default	weather-streaming-to-bronze-driver	1 (8%)	0 (0%)	1433Mi (11%)	1433Mi (11%)	3h30m
elastic-system	elastic-operator-0	100m (0%)	1 (8%)	150Mi (1%)	1Gi (8%)	9d
kube-system	coredns-66bc5c9577-4nllb	100m (0%)	0 (0%)	70Mi (0%)	170Mi (1%)	9d
kube-system	coredns-66bc5c9577-lx5d5	100m (0%)	0 (0%)	70Mi (0%)	170Mi (1%)	9d
kube-system	etcd-data-platform-control-plane	100m (0%)	0 (0%)	100Mi (0%)	0 (0%)	9d
kube-system	kindnet-hj7ln	100m (0%)	100m (0%)	50Mi (0%)	50Mi (0%)	9d
kube-system	kube-apiserver-data-platform-control-plane	250m (2%)	0 (0%)	0 (0%)	0 (0%)	9d
kube-system	kube-controller-manager-data-platform-control-plane	200m (1%)	0 (0%)	0 (0%)	0 (0%)	9d
kube-system	kube-proxy-zcgtg	0 (0%)	0 (0%)	0 (0%)	0 (0%)	9d
kube-system	kube-scheduler-data-platform-control-plane	100m (0%)	0 (0%)	0 (0%)	0 (0%)	9d
local-path-storage	local-path-provisioner-7b8c8ddb6-54m6x	0 (0%)	0 (0%)	0 (0%)	0 (0%)	9d
observability	es-weather-es-default-0	250m (2%)	1 (8%)	1Gi (8%)	1Gi (8%)	9d
observability	kibana-weather-kb-76fd7567b8-f4xv6	200m (1%)	1 (8%)	512Mi (4%)	1Gi (8%)	8d
spark-operator	spark-operator-controller-5fbfcb9d47-wshds	0 (0%)	0 (0%)	0 (0%)	0 (0%)	9d
spark-operator	spark-operator-webhook-55c6f74b59-xn4vm	0 (0%)	0 (0%)	0 (0%)	0 (0%)	9d

```
(big-data) minh@MacBook-Pro-cua-Minh kubernetes % kubectl get ns -o name
namespace/data
namespace/default
namespace/elastic-system
namespace/kafka
namespace/kube-node-lease
namespace/kube-public
namespace/kube-system
namespace/local-path-storage
namespace/observability
namespace/spark-operator
```

```
Allocated resources:
(Total limits may be over 100 percent, i.e., overcommitted.)
Resource           Requests           Limits
-----
cpu                 7700m (64%)       4100m (34%)
memory              10958Mi (91%)     12444Mi (104%)
ephemeral-storage   0 (0%)            0 (0%)
hugepages-1Gi       0 (0%)            0 (0%)
hugepages-2Mi       0 (0%)            0 (0%)
hugepages-32Mi      0 (0%)            0 (0%)
hugepages-64Ki      0 (0%)            0 (0%)
Events:              <none>
```


IV. Lesson learned

System Integration

- Enable data flow from **Kafka to Spark** and from **Spark to Elasticsearch**.
- Kafka and Spark run on **Kubernetes**, requiring proper service configuration for communication.

Approaches Tried

- **Docker Compose Configuration**
 - Used docker-compose.yml to connect Kafka, Spark, and Elasticsearch.
 - Effective for local development and testing.
- **Kubernetes YAML Configuration**
 - Used .yaml manifests to define services, deployments, and networking.
 - Kubernetes-native approach for inter-service communication.

Final Solution: Kubernetes YAML files selected

- Required for Kubernetes-based deployment.
- Provides full control over networking, scaling, and service discovery.

IV. Lesson learned

Performance Optimization

- **Problem:** Frequent loading of S3 data for each ML inference.
- **Approach:**
 - Cache historical data in Spark DataFrames.
 - Load only the previous 24 hours at the start of each hour.
- **Final Solution:** Use cached data instead of reloading from S3.

Monitoring & Debugging

Problem: Need system logs and visibility for monitoring.

Approaches

- Kubernetes CLI commands (kubectl describe, logs).
- Grafana-based monitoring.

Final Solution

- Kubernetes-native commands used.
- Grafana rejected due to high resource overhead.



IV. Lesson learned

Scaling

- **Problem:** Improve Spark processing performance.
- **Approach:**
 - Increased Spark executors from 1 to 2.
 - Enabled parallel processing of data partitions.
- **Final Solution:** Two Spark executors used in production.



Data Quality & Testing

- **Problem:** Prevent broken pipelines before pushing code to GitHub.
- **Approaches**
 - Local unit testing of individual components.
 - Manual testing of data flow.
 - Docker used to simulate production services.
- **Final Solution:** Combine local testing with container-based integration testing.

IV. Lesson learned

Data Quality & Testing

- **Problem**
 - Prevent broken pipelines before pushing code to GitHub.
- **Approaches**
 - Local unit testing of individual components.
 - Manual testing of data flow.
 - Docker used to simulate production services.
- **Final Solution**
 - Combine local testing with container-based integration testing.
- **Key Takeaways**
 - Early testing reduces deployment errors.
 - Component isolation simplifies debugging.

V. Conclusion

In conclusion, we have ...

- This project successfully builds an end-to-end real-time weather data pipeline for Vietnam using modern Big Data technologies.
- The system integrates Kafka, Spark (streaming & batch), Elasticsearch, and Kubernetes to process, store, and analyze high-velocity weather data.
- A Lambda Architecture is applied to support both real-time analytics and historical batch processing.
- The pipeline demonstrates scalability, fault tolerance, and modular deployment suitable for real-world data engineering systems.

Future works

- Extend forecasting models with advanced ML/DL approaches (e.g., LSTM, Transformer-based time-series models).
- Expand data sources to include additional meteorological APIs and sensor networks.
- Deploy the system on a multi-node Kubernetes cluster to evaluate scalability and high availability.
- Optimize resource usage and monitoring with lightweight observability solutions.